

car price prediction

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('CarPrice.csv')
```

```
In [3]: df.head()
```

Out[3]:

| | car_ID | symboling | CarName | fueltype | aspiration | doornumber | carbody | drivewheel | engine | location | wt |
|---|--------|-----------|-----------------------------|----------|------------|------------|-------------|------------|--------|----------|----|
| 0 | 1 | 3 | alfa-romero giulia | gas | std | two | convertible | rwd | | front | |
| 1 | 2 | 3 | alfa-romero stelvio | gas | std | two | convertible | rwd | | front | |
| 2 | 3 | 1 | alfa-romero Quadrifoglio | gas | std | two | hatchback | rwd | | front | |
| 3 | 4 | 2 | audi 100 ls | gas | std | four | sedan | fwd | | front | |
| 4 | 5 | 2 | audi 100ls | gas | std | four | sedan | 4wd | | front | |

5 rows × 26 columns



```
In [4]: df.tail()
```

Out[4]:

| | car_ID | symboling | CarName | fueltype | aspiration | doornumber | carbody | drivewheel | engine | location | whe |
|-----|--------|-----------|--------------------|----------|------------|------------|---------|------------|--------|----------|-----|
| 200 | 201 | -1 | volvo 145e (sw) | gas | std | four | sedan | rwd | | front | |
| 201 | 202 | -1 | volvo 144ea | gas | turbo | four | sedan | rwd | | front | |
| 202 | 203 | -1 | volvo 244dl | gas | std | four | sedan | rwd | | front | |
| 203 | 204 | -1 | volvo 246 | diesel | turbo | four | sedan | rwd | | front | |
| 204 | 205 | -1 | volvo 264gl | gas | turbo | four | sedan | rwd | | front | |

5 rows × 26 columns



In [5]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName                205 non-null   object
3   fueltype              205 non-null   object
4   aspiration             205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

In [6]: df.describe()

Out[6]:

| | car_ID | symboling | wheelbase | carlength | carwidth | carheight | curbweight | enginesize | bo |
|--------------|------------|------------|------------|------------|------------|------------|-------------|------------|------------|
| count | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 |
| mean | 103.000000 | 0.834146 | 98.756585 | 174.049268 | 65.907805 | 53.724878 | 2555.565854 | 126.907317 | 3.5 |
| std | 59.322565 | 1.245307 | 6.021776 | 12.337289 | 2.145204 | 2.443522 | 520.680204 | 41.642693 | 0.2 |
| min | 1.000000 | -2.000000 | 86.600000 | 141.100000 | 60.300000 | 47.800000 | 1488.000000 | 61.000000 | 2.5 |
| 25% | 52.000000 | 0.000000 | 94.500000 | 166.300000 | 64.100000 | 52.000000 | 2145.000000 | 97.000000 | 3.1 |
| 50% | 103.000000 | 1.000000 | 97.000000 | 173.200000 | 65.500000 | 54.100000 | 2414.000000 | 120.000000 | 3.5 |
| 75% | 154.000000 | 2.000000 | 102.400000 | 183.100000 | 66.900000 | 55.500000 | 2935.000000 | 141.000000 | 3.5 |
| max | 205.000000 | 3.000000 | 120.900000 | 208.100000 | 72.300000 | 59.800000 | 4066.000000 | 326.000000 | 3.5 |

```
In [7]: df.isnull().sum()
```

```
Out[7]: car_ID          0
        symboling      0
        CarName        0
        fueltype       0
        aspiration     0
        doornumber     0
        carbody        0
        drivewheel     0
        enginelocation 0
        wheelbase      0
        carlength      0
        carwidth       0
        carheight      0
        curbweight     0
        enginetype     0
        cylindernumber 0
        enginesize     0
        fuelsystem     0
        boreratio      0
        stroke         0
        compressionratio 0
        horsepower     0
        peakrpm        0
        citympg        0
        highwaympg     0
        price          0
        dtype: int64
```

```
In [8]: df.duplicated().sum()
```

```
Out[8]: 0
```

```
In [9]: df.shape
```

```
Out[9]: (205, 26)
```

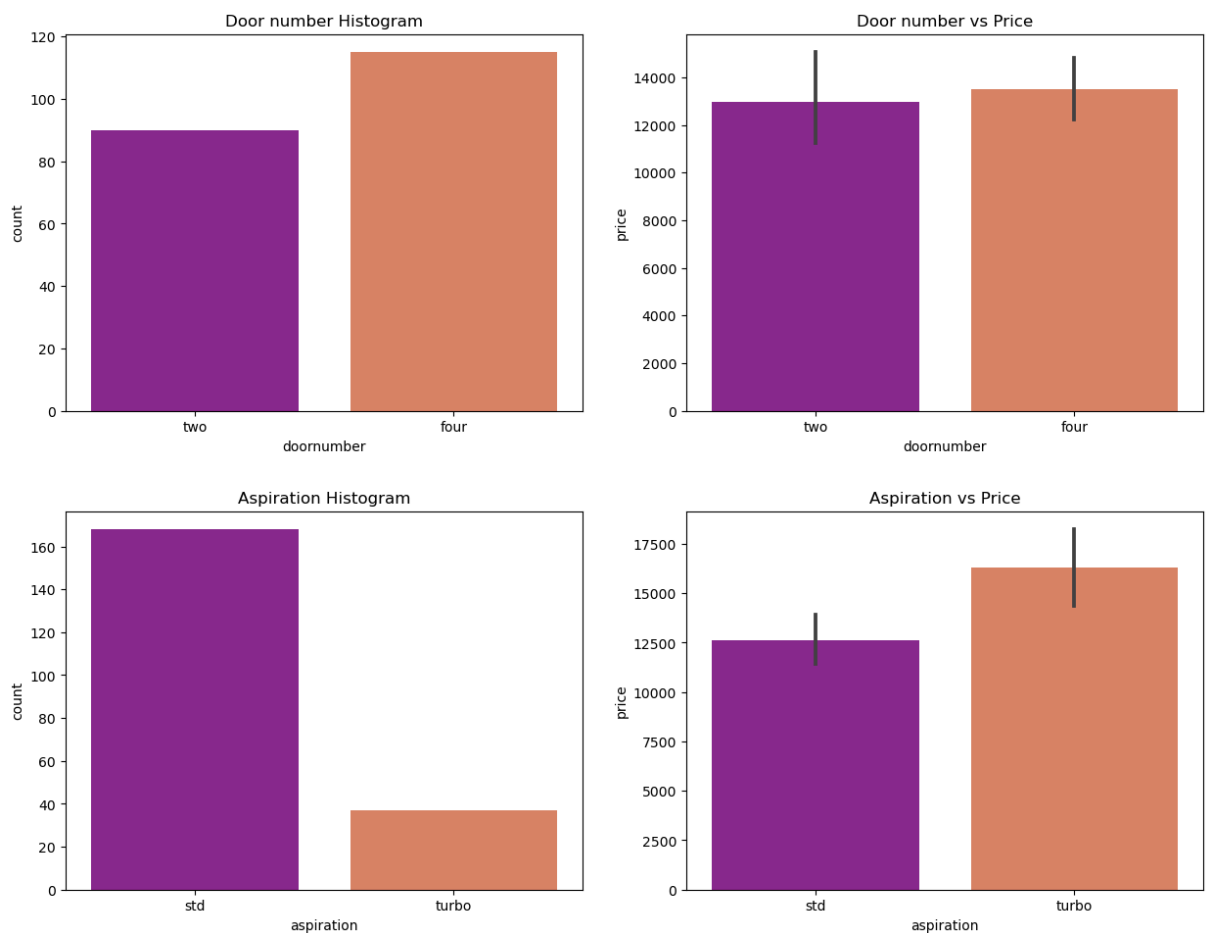
```
In [10]: print(df.price.describe(percentiles=[0.225,0.50,0.75,0.85,0.98,1]))
```

```
count      205.000000
mean       13276.710571
std         7988.852332
min         5118.000000
22.5%       7609.000000
50%        10295.000000
75%        16503.000000
85%        18500.000000
98%        36809.600000
100%       45400.000000
max        45400.000000
Name: price, dtype: float64
```

Data Visualization

```
In [11]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.title("Door number Histogram")
sns.countplot(data=df, x='doornumber', palette="plasma")
plt.subplot(1,2,2)
plt.title('Door number vs Price')
sns.barplot(data=df, x='doornumber', y='price', palette="plasma")
plt.show()


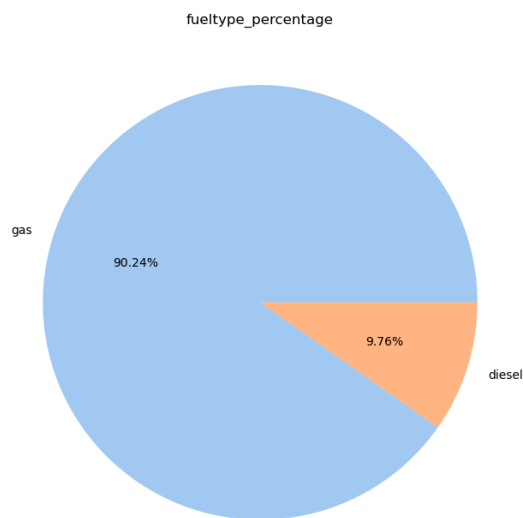
plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.title("Aspiration Histogram")
sns.countplot(data=df, x='aspiration', palette="plasma")
plt.subplot(1,2,2)
plt.title("Aspiration vs Price")
sns.barplot(data=df, x='aspiration', y='price', palette="plasma")
plt.show()
```



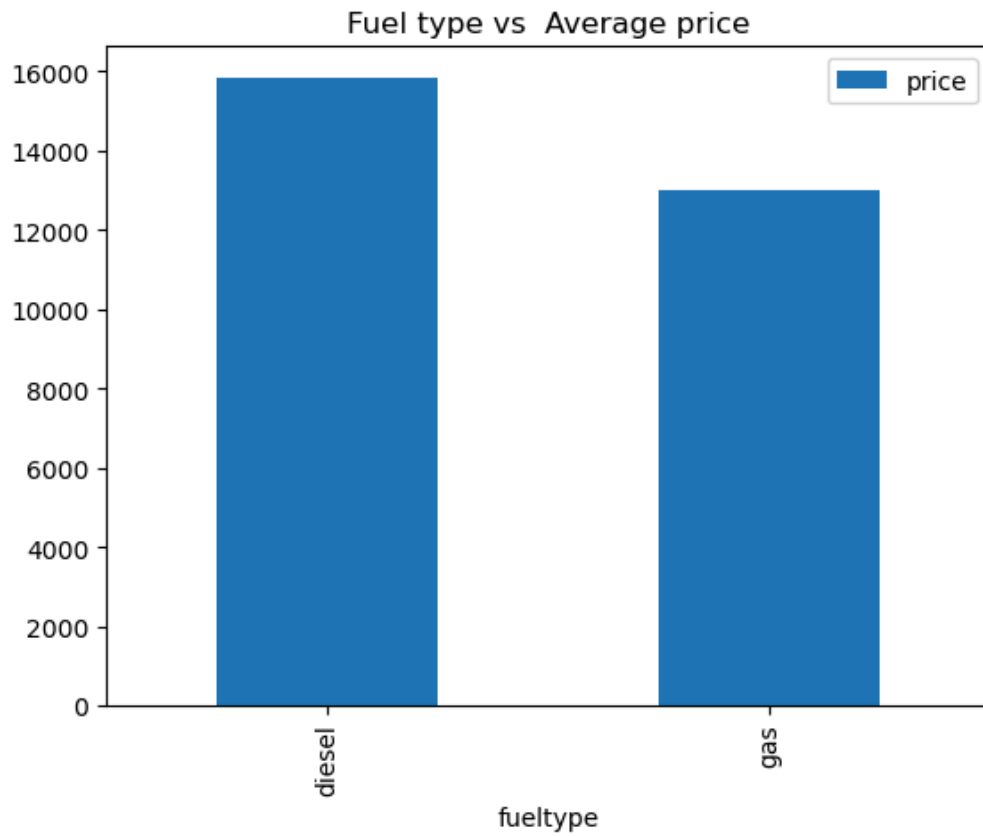
```
In [12]: colors=sns.color_palette('pastel')
labels=df['fueltype'].dropna().unique()
plt.figure(figsize=(18,10))
plt.subplot(1,2,1)

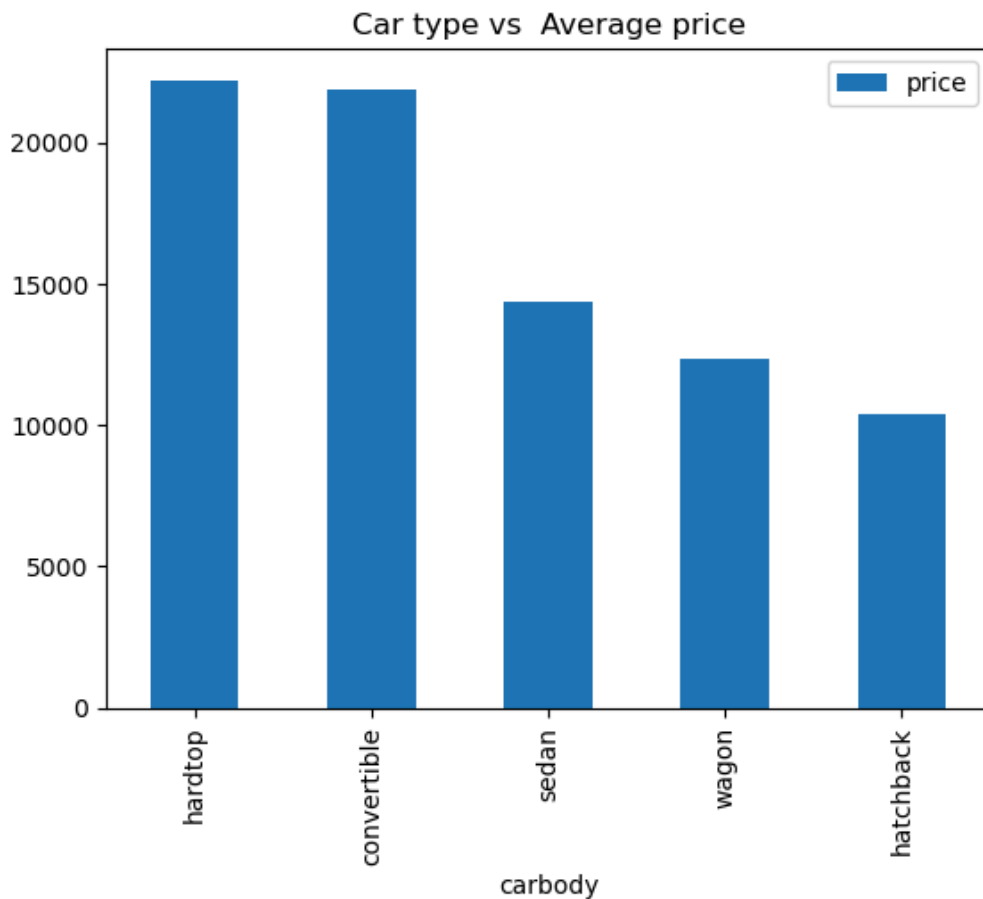
plt.title('fueltype_percentage')
plt.pie(df['fueltype'].value_counts(),labels=labels,colors=colors,autopct='%.2f%%')
plt.subplot(1,2,2)
plt.title('fueltype bar chart')
sns.countplot(x='fueltype',data=df,palette=colors)
df.fueltype.value_counts(dropna=False)
```

```
Out[12]: fueltype
gas      185
diesel   20
Name: count, dtype: int64
```



```
In [13]: dff=pd.DataFrame(df.groupby(['fueltype'])['price'].mean().sort_values(ascending=False))
dff.plot.bar()
plt.title("Fuel type vs Average price")
plt.show()
dff=pd.DataFrame(df.groupby(['carbody'])['price'].mean().sort_values(ascending=False))
dff.plot.bar()
plt.title("Car type vs Average price")
plt.show()
```





```
In [15]: y=df['price']
x=df[['symboling','wheelbase','carwidth','carheight','curbweight','enginesize','borera
```

```
In [16]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=100
```

Random Forest Model

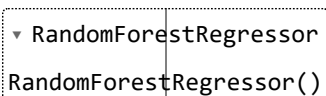
```
In [17]: from sklearn.ensemble import RandomForestRegressor
```

```
In [18]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=100
print('training data shape is:{}.'.format(x_train.shape))
print('training label shape is:{}.'.format(y_train.shape))
print('testing data shape is:{}.'.format(x_test.shape))
print('testing data shape is:{}.'.format(y_test.shape))
```

```
training data shape is:(164, 13).
training label shape is:(164,).
testing data shape is:(41, 13).
testing data shape is:(41,).
```

```
In [19]: from sklearn.ensemble import RandomForestRegressor
regressor=RandomForestRegressor()
```

In [20]: regressor.fit(x,y)

Out[20]: 

In [21]: regressor.score(x_train,y_train)

Out[21]: 0.9883097125128573

In [22]: regressor.score(x_test,y_test)

Out[22]: 0.9867846662115134

In [23]: `from sklearn.metrics import accuracy_score`
`predictions=regressor.predict(x_test)`

In [24]: `percentage=regressor.score(x_test,y_test)`
`percentage`

Out[24]: 0.9867846662115134

In [25]: `print(regressor.score(x_train,y_train))`
`print(f"test set:{len(x_test)}")`
`print(f"Accuracy={percentage*100}%")`

0.9883097125128573
test set:41
Accuracy=98.67846662115134%

Thank you

In []: